

Architecting Big Data Ecosystems

An overview of NOSQL Systems and Apache Spark

Dr Asma ZGOLLI

Senior Machine Learning Engineer

Guest Lecturer at SRH University of Applied Sciences
Hamburg

Invited by:

Prof Dr Nnamdi Oguji

8 November 2023

Big data ecosystem

Introduction

➤ Applications:

- Demands and sales forecast
- Predictive maintenance
- Fraud and anomaly detection
- Efficient data management

➤ Data:

- Orders and payment
- Product catalogues and clients profile
- Smart meters, network, and sensors
- Weather and geospatial information

✓ High volume, variety, veracity, and velocity

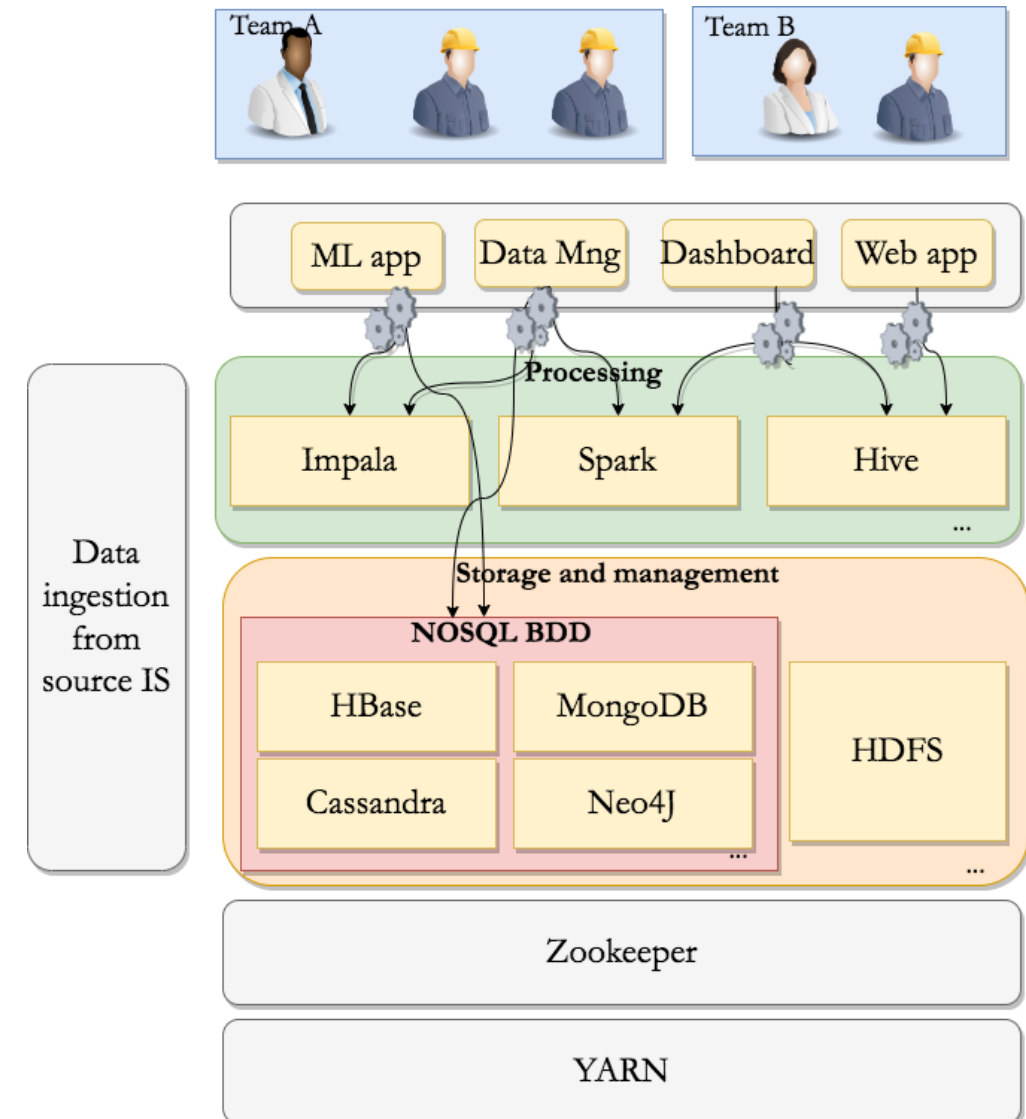


Figure 1: Example of a Big Data Architecture

Roadmap

Big Data Ecosystems

- Introduction

NOSQL Datastores

- NOSQL
- NOSQL Data Models
- Key-Value Datastores
- Document Datastores
- Column-Oriented Datastores
- Graph Datastores
- Comparing NoSQL Data Models: A Snapshot

Apache Spark

- Spark Internals
- Spark APIs

- NOSQL: **Not Only SQL**
- **Scalability:** *horizontal scalability on commodity hardware using distribution techniques.*
- **Performance vs Transactions:** NOSQL prioritizes performance over transactional semantics.
- **Complex Queries:** *In addition to key-only access, NOSQL supports richer queries e.g.*
 - Multi-property indexing and sophisticated data retrieval. (MongoDB)
 - Column-based filtering. (BigTable)
- **Big Data Integration**
 - Bindings to Hadoop or other parallel and distributed processing frameworks, such as MapReduce.
- **Schema-Free Storage**
 - *Supports semi/unstructured data*
 - (+) Faster application development and less costly on-the-fly schema updates
 - (-) More responsibility on application developers and the need for versioning

NOSQL Data Models

Data model	Data structure	Relation	Record	Attribute	Data view
Document	Hierarchical tree e.g. : map , collection	Collection	Document	Attribute	Physical view
Key-value	Blob , text , json , xml	Bucket	Key-value pair	X	Physical view
Column family	Tabular data	Column family	Row	Column	Logical view
Graph	Graph , tree	Node	Node	Attribute	Logical view
		Link	link	Attribute	

Key-Value Datastores ...

- **Physical Data Model:** $\langle \text{key}, \text{value} \rangle$ pairs.
- **Key**
 - Hash key
- **Value**
 - Blob
 - Complex, non-typed: Text, structured, multivalued (array), or atomic
- **Key-Based Data Access:**
 - Supported operators:
 - `put(key, value);`
 - `value= get(key);`
 - `delete(key);`
 - Very efficient data access.
 - Limited API
 - Increased complexity related to data processing at the application level.
- **Use Cases :** Sessions management, data indexing

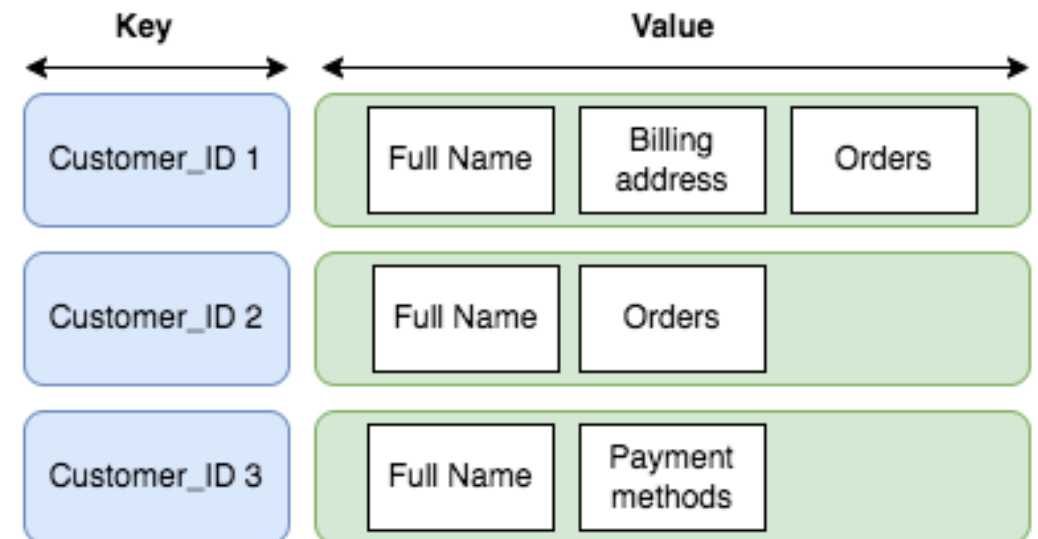


Figure 2: Example of a key value pairs collection

... Key-Value Datastores

Case study: Dynamo

- A pure key-value store [G23,DHJK+07]
 - Scalable
 - Decentralized (Peer-to-peer architecture)
 - Highly available
- Data partitioning & Replication:
 - Consistent hashing : Keys are stored at the next node in the ring with a higher ID.
 - Virtual Nodes: Serve as an abstract physical nodes for enhanced fault tolerance and load balancing.
- Amazon Use Cases:
 - Powers critical services: shopping carts, bestseller lists, sales rankings.
 - Distinct from DynamoDB, which is offered by AWS.

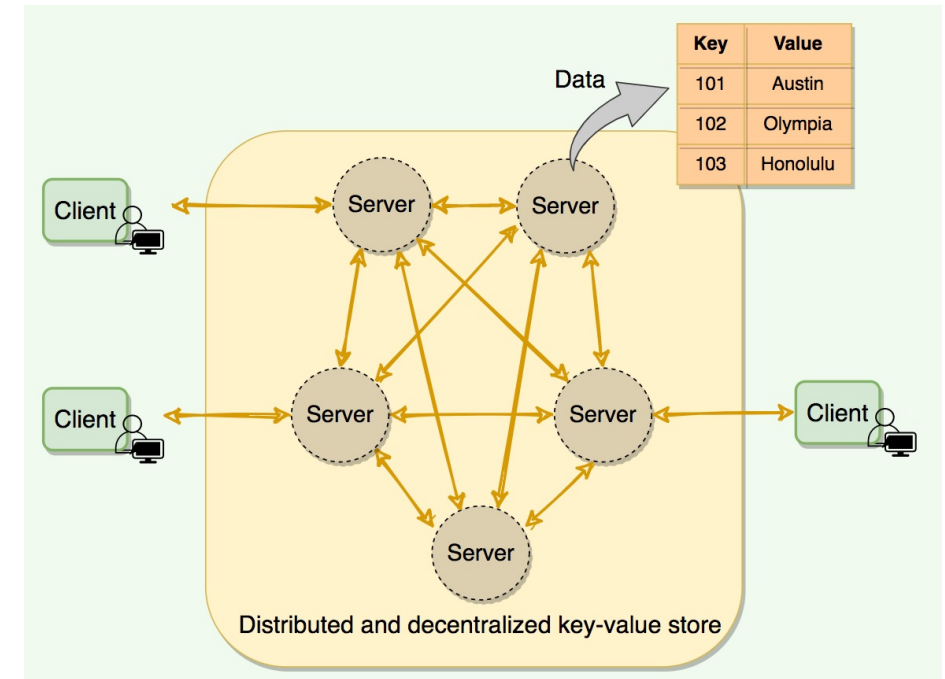


Figure 3: key-value data store distribution model [G23]




Figure 4: Consistent Hashing

Column-Oriented Datastores ...

- **Physical Data Model:**
 - Data is stored by column
- **Logical Data Model:**
 - Big table
- **Column-Family:**
 - A group of columns stored together in the same server optimizing data retrieval
- **Column:**
 - Basic data storage unit
 - atomic (single data type) or structured (nested data).
- **Column-Oriented Data access:**
 - Reads only the **relevant columns** which is suitable for read-intensive data access.
 - Writing data requires significant **IO**.
 - Supports operations such as filtering and aggregation, but does **not support joins** or **subqueries**.

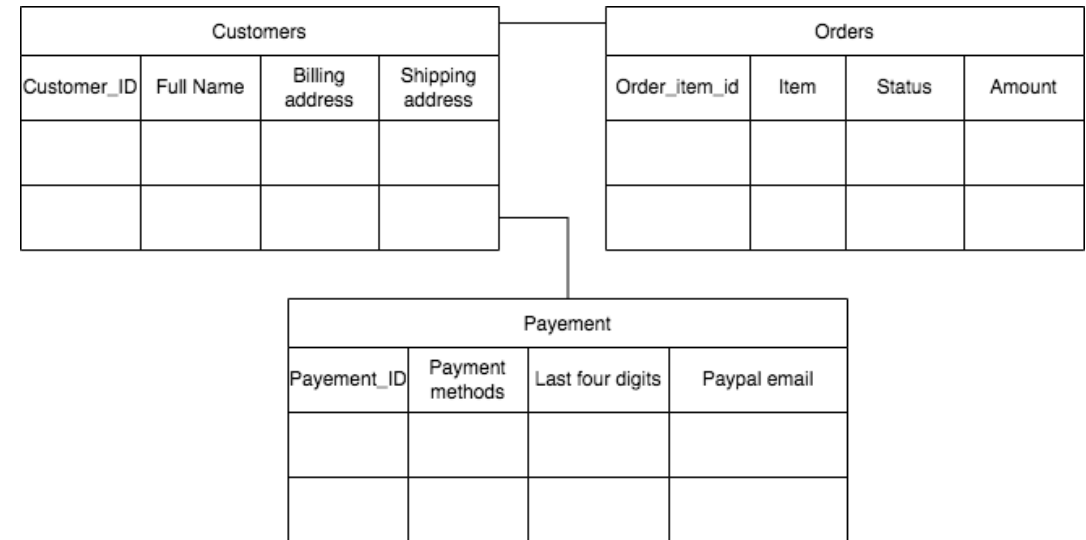


Figure 5: Example of tables (Relational data model)

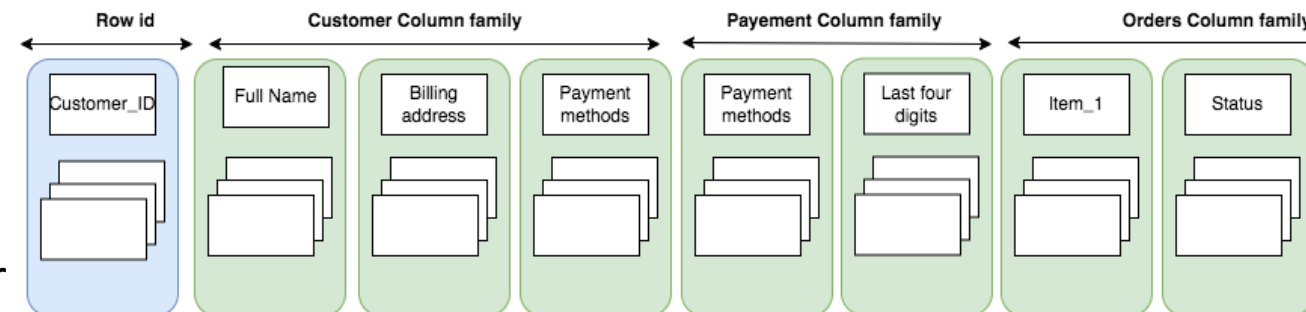


Figure 6: Example of a column-family

Column-Oriented Datastores ...

NOSQL Datastores

Case study: Hbase

- Hybrid Data Model :
 - key-value + column oriented.
- Foundation: Built on top of the Hadoop Distributed File System (HDFS).
- Distributed Architecture:
 - Replication : e.g. master-slave, p2p clusters, etc.
 - Partitioning : auto-sharding.
 - Parallelism : Write-Ahead Logging (WAL) and Bloom filters.
- System components [Tea16]:
 - HMaster: Load balancing, catalogue management, and orchestration.
 - RegionServer: Manages data partitions (Regions)
 - HFile: Physical storage format.
 - MemStore: In-memory write buffer.

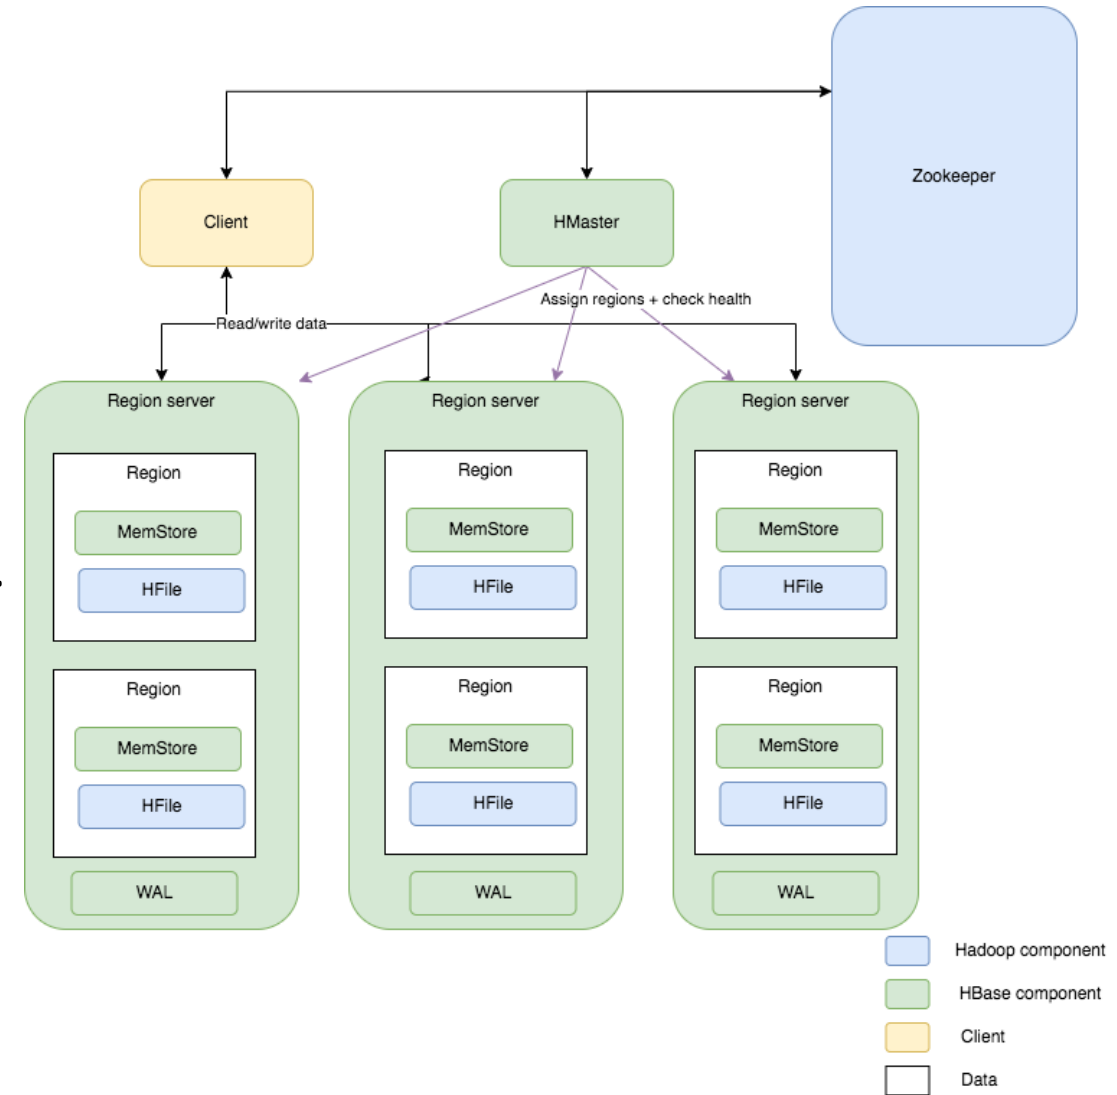


Figure 7: Apache Hbase Architecture

- **Physical Data Model:**
 - JSON like
 - Flexible schema
 - Hierarchical
 - Stores of complex data in fields called **aggregates**.
- **Rich API:**
 - Advanced data access & filtering
 - Built-in aggregation functions
 - No joins
- **Use cases:**
 - Products catalogues
 - Clients profiles
 - Content management systems (e.g. CRM)

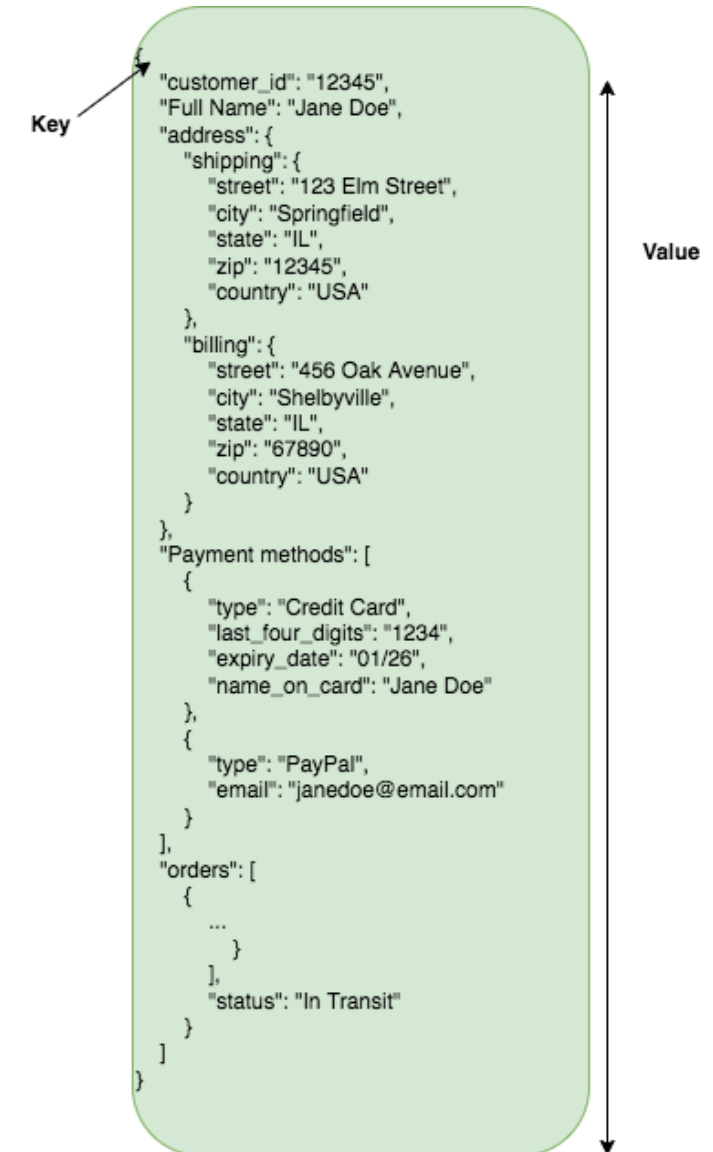


Figure 8: Example of a document

... Document-Oriented Datastores

Case study: MongoDB

- Data model:
 - BSON (Binary JSON)
- Documents:
 - Typed fields
 - Arrays and nested documents
- Data partitioning & Replication:
 - Master / slave replication
 - Replica set
 - Automatic failover
 - Partitioning : sharding
- Indexing:
 - Utilizes a B-tree data structure.

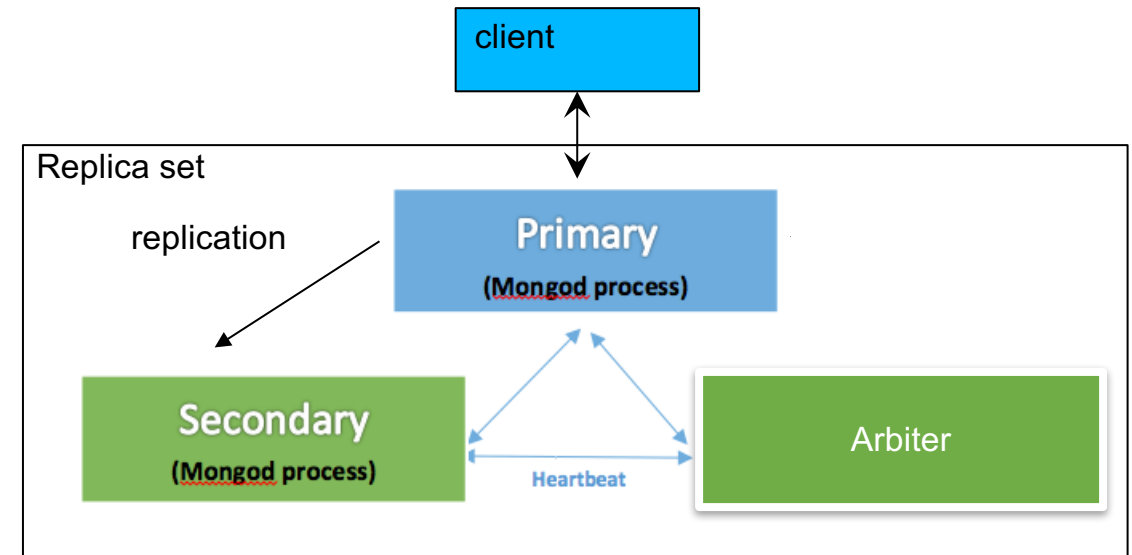


Figure 9: MongoDB replication architecture




Figure 10: MongoDB partitioning architecture [Mon23]

- Graphs are data structures that help modelling complex relationships between different entities (real-world objects).

- **Nodes:**

- Also called vertices
- Represents the entities (real-world objects)

- **Links:**

- Also called edges
- represents the relationship between these nodes

- **Physical data model:**

- Nodes and links typically stored separately.

- **Use cases:**

- Search and pathfinding in a graph ,
recommender system, social media data
management

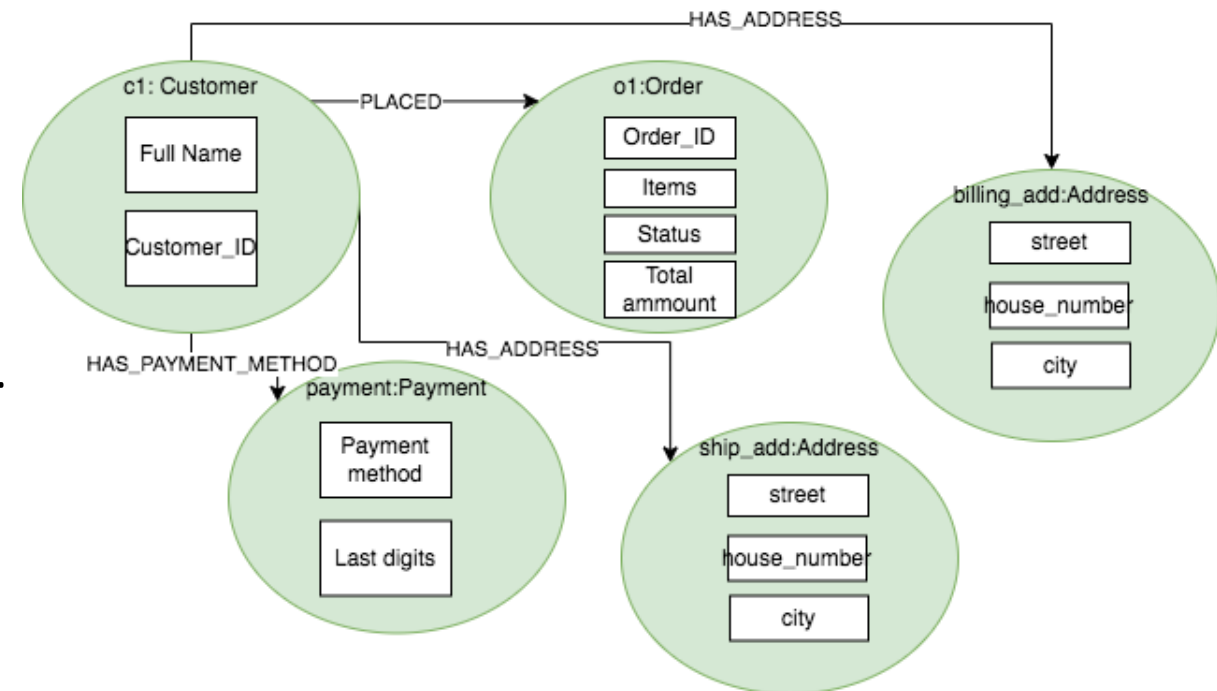


Figure 11: Example of a graph

... Graph Datastores

Case study: Tigergraph

- Data model:
 - Vertices and Edges with attribute support.
 - Attributes can be atomic or structured.
- GSQL:
 - Graph-specific, SQL-like declarative language.
 - Supports graph pattern matching.
- Scalability:
 - Horizontally scalable architecture.
- Processing:
 - Massively Parallel Processing (MPP).

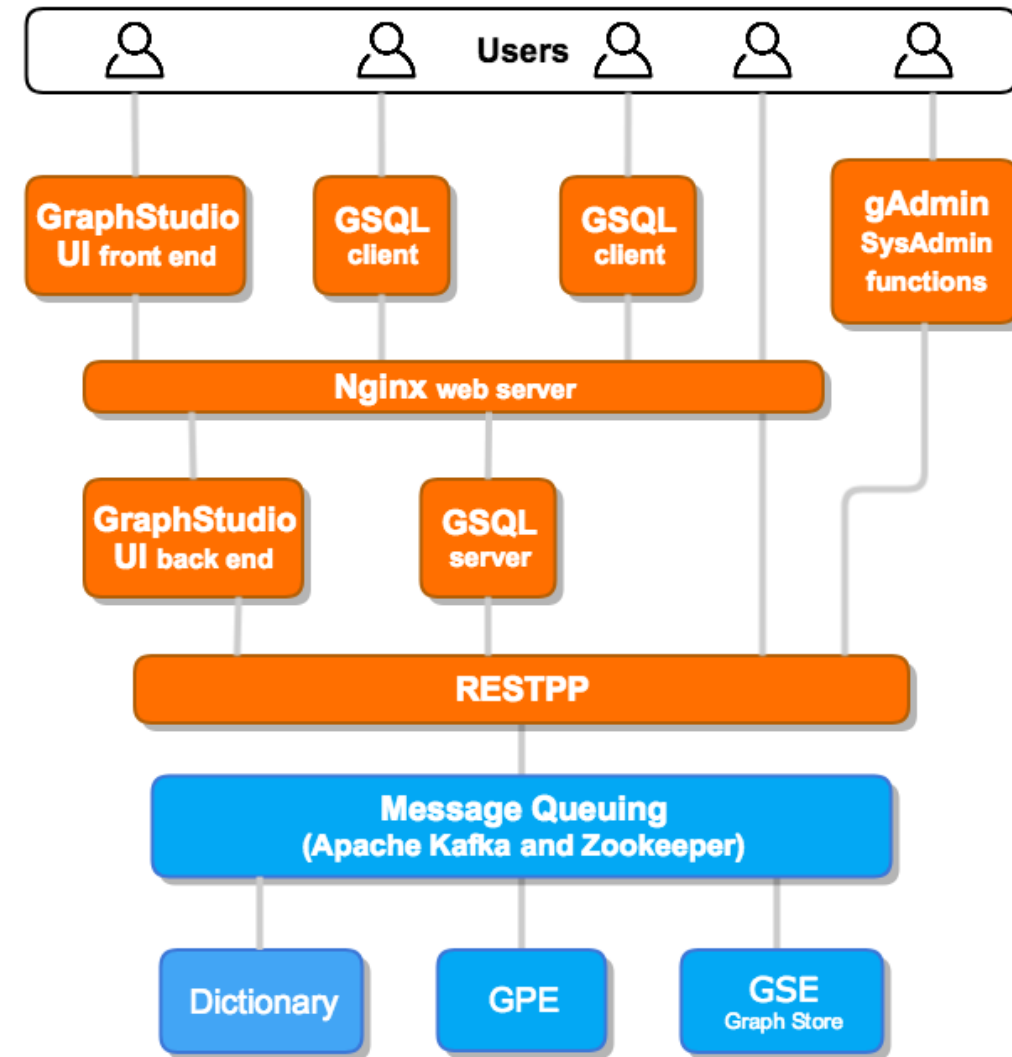


Figure 12: Tigergraph architecture [Tig23]

Comparing NoSQL Data Models: A Snapshot [ZBV21] ...

Data model	Principle	Advantage	Disadvantage
Key-value	Data stored and accessed as a key-value pair: Key a unique identifier and Complex data as value	<ul style="list-style-type: none">✓ Efficient index on data✓ Simple data model	Limited access API
Document oriented	Aggregates stored as documents	<ul style="list-style-type: none">✓ Complex queries✓ Rich filters	Slow writes, lack of consistency
Column family oriented	Data organised as columns. Columns store aggregates and are organised as column families	<ul style="list-style-type: none">✓ Evolving data schema✓ Efficient column oriented storage	Slow writes, bad performance with small data and row based access
Graph oriented	Representing the connexions between data records	<ul style="list-style-type: none">✓ Graph algorithms for efficient connected data querying	Does not support well distribution

... Comparing NoSQL Data Models: A Snapshot [ZBV21]

Operator / data model	Key value (e.g Dynamo)	Document based (e.g MongoDB)	Column family (e.g Apache HBase)	Graph based (e.g Neo4J)
Scan	XXX	X	XX	X
Filter	X	XX	X	XX
Join	-	X	-	XXX
Aggregate	-	XXX	X	XX
Group by	-	XXX	X	XX
Order by	-	X	X	X

The number of x denotes the performance of the operator (x, xx, xxx)

Roadmap

Big Data Ecosystems

- Introduction

NOSQL Datastores

- NOSQL
- NOSQL Data Models
- Key-Value Datastores
- Document Datastores
- Column-Oriented Datastores
- Graph Datastores
- Comparing NoSQL Data Models: A Snapshot

Apache Spark

- Spark Internals
- Spark APIs

➤ Apache Spark

- **Cluster computing framework integrated** in the Hadoop ecosystem:
 - Horizontally scalable data processing
 - Parallel data processing
- Characteristics :
 - **In memory processing**: it addresses the limitation of MapReduce that stores intermediate results on disk after every task by using a caching technique
 - **Distributed processing**: a **master** node receives the workload from the driver program and coordinates the processing by sends it to the **worker** nodes for execution.
 - **Lazy evaluation**: tasks are classified as actions and transformations and structured as a **DAG**. The processing is only triggered when an action is called.




Figure 13: Apache Spark Architecture [Spa]

➤ Spark RDD API

- RDDs are a collection of immutable (in-memory read-only) objects that are distributed across the cluster
- Low level API: RDDs are the physical data model of Spark, other APIs are built on top of it to facilitate the development in spark

• Example:

Read the input file and Calculating words count

```
text_file = sc.textFile("firstprogram.txt")
counts = text_file.flatMap(lambda line: line.split(" ")) \
                    .map(lambda word: (word, 1)) \
                    .reduceByKey(lambda x, y: x + y)
```

Printing each word with its respective count

```
output = counts.collect()
for (word, count) in output:
    print("%s: %i" % (word, count))
```

➤ **Spark DataFrame**

- Collection of immutable and distributed data with named columns
- Similar data model to a relational table
- Has a Domain Specific Language API that allows querying and transforming the dataframes

➤ **SparkSQL**

- Interface that allows executing an SQL-like query language on dataframes
- Supports query optimization (Catalyst optimizer)

➤ **Data Science libraries**

- MLlib: scalable and simple ML library that is integrated with spark APIs
- Graphx: scalable and parallel processing for graphs

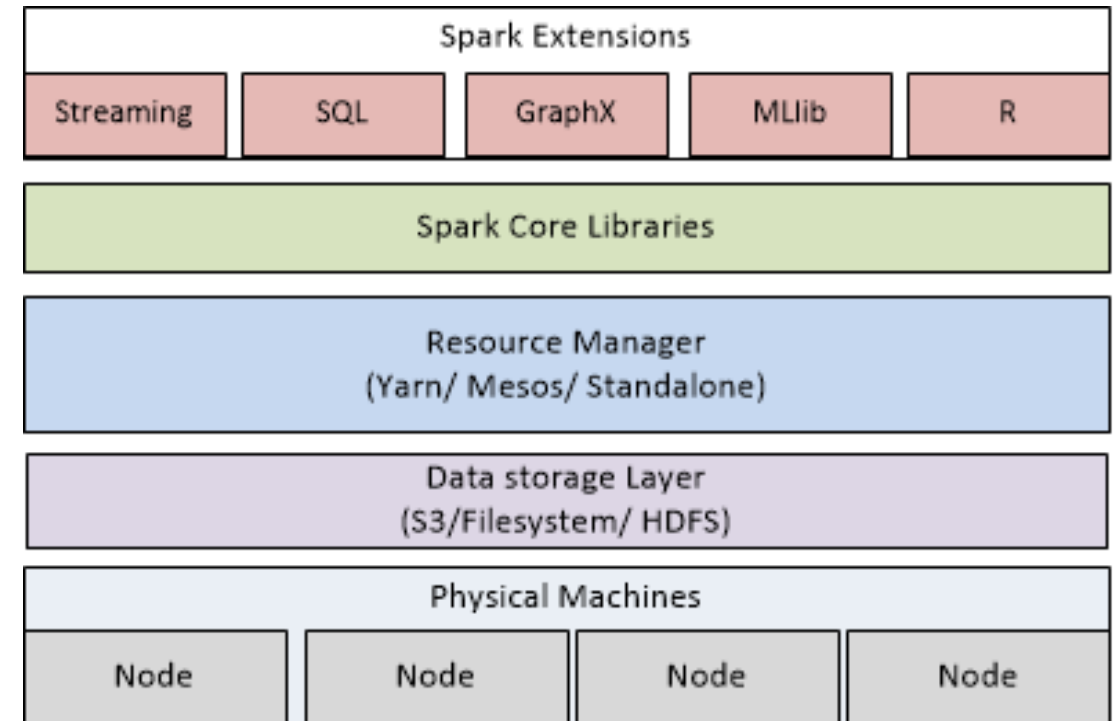


Figure 14: Apache Spark Components [SG17]

Thank you for your attention



Bibliography...

- [G23] Design Gurus. (n.d.). Dynamo: Introduction. Design Gurus: One-Stop Portal For Tech Interviews. Retrieved November 2, 2023, from <https://www.designgurus.io/course-play/grokking-the-advanced-system-design-interview/doc/636ce5b668ee620a5aa631d2>
- [DHJK+07] Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Voshall, and Werner Vogels. 2007. Dynamo: amazon's highly available key-value store. SIGOPS Oper. Syst. Rev. 41, 6 (December 2007), 205–220. <https://doi.org/10.1145/1323293.1294281>
- [Mon23] MongoDB. MongoDB documentation. Available at <https://docs.mongodb.com>, Accessed: 2023-10-10, 2023.
- [SF12] Pramod J. Sadalage and Martin Fowler. 2012. NoSQL Distilled: A Brief Guide to the Emerging World of Polyglot Persistence (1st. ed.). Addison-Wesley Professional.
- [SG17] Saxena, Shilpi, and Saurabh Gupta. 2017. Practical Real-time Data Processing and Analytics. Packt Publishing.

... Bibliography

- [Spa]Team Spark. Cluster Mode Overview - Spark 3.1.1 Documentation. Available at <https://spark.apache.org/docs/latest/cluster-overview.html>, Accessed: 2023-10-10.
- [Tea16] Apache HBase Team. Hbase™ reference guide. Available at <https://hbase.apache.org/1.2/book.html>, Accessed: 2023-10-10, 2016.
- [Tig23] TigerGraph. TigerGraph docs. Available at <https://docs.tigergraph.com/home/>, Accessed: 2023-10-10, 2023.
- [ZBV21] Asma Zgolli, Christophe Bobineau, Genoveva Vargas-Solar. 2021. Recommandation de placement de données pour les traitements dans des lacs de données smart grids. [Data placement recommendation for processing in smart grid data lakes]. (Thesis Presentation).